




# Chapter 7: Transactions





# Contents

- Transaction Concept
- Transaction State
- Concurrent Executions
- Serializability
- Recoverability
- Implementation of Isolation
- Transaction Definition in SQL
- Testing for Serializability.



# Transaction Concept

- Collection of operations that form a single logical unit of work are called **transactions**

**Example**

Transaction to transfer \$50 from account A to account B:

1.  $\text{read}(A)$
2.  $A := A - 50$
3.  $\text{write}(A)$
4.  $\text{read}(B)$
5.  $B := B + 50$
6.  $\text{write}(B)$

Two main issues to deal with:

1. Failures of various kinds, such as hardware failures and system crashes
2. Concurrent execution of multiple transactions



# Example of Fund Transfer<sup>Cont.</sup>

Transaction to transfer \$50 from account A to account B:

1. `read(A)`
2. `A := A - 50`
3. `write(A)`
4. `read(B)`
5. `B := B + 50`
6. `write(B)`

System  
Failure

- **Atomicity requirement**
  - if the transaction fails after step 3 and before step 6, money will be “lost” leading to an inconsistent database state
    - Failure could be due to software or hardware
  - the system should ensure that updates of a partially executed transaction are not reflected in the database



# Example of Fund Transfer <sup>Cont.</sup>



Transaction to transfer \$50 from account A to account B:

1. `read(A)`
2. `A := A - 50`
3. `write(A)`
4. `read(B)`
5. `B := B + 50`
6. `write(B)`

- **Durability requirement**
  - once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place), the updates to the database by the transaction must persist even if there are software or hardware failures.



# Example of Fund Transfer<sup>Cont.</sup>

Transaction to transfer \$50 from account A to account B:

1. read( $A$ )
2.  $A := A - 50$
3. write( $A$ )
4. read( $B$ )
5.  $B := B + 50$
6. write( $B$ )


$$A+B = 500$$

$$A+B = 450$$

$$A+B = 500$$

- **Consistency requirement**
  - the sum of A and B is unchanged by the execution of the transaction

- A transaction must see a consistent database.
- During transaction execution the database may be temporarily inconsistent.
- When the transaction completes successfully the database must be consistent





# Example of Fund Transfer <sup>Cont.</sup>

Transaction  $T_1$  to transfer \$50 from account A to account B, Transaction  $T_2$  print out the sum of the balance of A and B

- Isolation requirement

- if between steps 3 and 6, another transaction  $T_2$  is allowed to access the partially updated database, it will see an inconsistent database (the sum  $A + B$  will be less than it should be).

$T_1$

1. read( $A$ )
2.  $A := A - 50$
3. write( $A$ )

4. read( $B$ )
5.  $B := B + 50$
6. write( $B$ )

$T_2$

1. read( $A$ )
2. read( $B$ )
3. print( $A+B$ )

Output  
 $A+B = 450$



# Example of Fund Transfer <sup>Cont.</sup>

Transaction  $T_1$  to transfer \$50 from account A to account B, Transaction  $T_2$  print out the sum of the balance of A and B


$T_1$

1. `read(A)`
2. `A := A - 50`
3. `write(A)`
4. `read(B)`
5. `B := B + 50`
6. `write(B)`

$T_2$

1. `read(A)`
2. `read(B)`
3. `print(A+B)`

- Isolation can be ensured trivially by running transactions **serially**
  - that is, one after the other.
- However, executing multiple transactions concurrently has significant benefits, as we will see later.








# ACID Properties

To preserve the integrity of data in the database system, a **transaction** must ensure:

- **Atomicity**. Either all operations of the transaction are properly reflected in the database or none are.
- **Consistency**. Execution of a transaction in isolation preserves the consistency of the database.




Cont.

# ACID Properties

- **Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions.

That is, for every pair of transactions  $T_i$  and  $T_j$ , it appears to  $T_i$  that either  $T_j$  finished execution before  $T_i$  started, or  $T_j$  started execution after  $T_i$  finished.



# ACID Properties <sup>Cont.</sup>

- **Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.



# Transaction State

- **Active** – the initial state; the transaction stays in this state while it is executing
- **Partially committed** – after the final statement has been executed.
- **Failed** -- after the discovery that normal execution can no longer proceed.
- **Committed** – after successful completion.



# Transaction State


Cont.

- **Aborted** – after the transaction has been rolled back and the database restored to its state prior to the start of the transaction. Two options after it has been aborted:
  - **restart the transaction**
    - can be done only if no internal logical error
  - **kill the transaction**



# Transaction State

Cont.






# Implementation of Atomicity and Durability

- The **recovery-management** component of a database system implements the support for atomicity and durability
  - Log file
  - Backup



# Log based recovery

- A **log** is kept on stable storage.
  - The log is a sequence of **log records**, and maintains a record of update activities on the database.
  - Before  $T_i$  executes  $\text{write}(X)$ , it write a log record



# Log based recovery

- When transaction  $T_i$  starts, it registers itself by writing a log record

$T_i$  start

- When  $T_i$  finishes its last statement, it writes a log record

$T_i$  commit



# Log based recovery



- Two approaches using logs
  - Deferred database modification
    - Can only output updates to hard disk while in partial commit state
    - Using Redo
  - Immediate database modification
    - Can output updates to hard disk while still in the active state
    - Using Undo, Redo



Output log record to hard disk first





# Checkpoints




- **Problems in recovery procedure as discussed earlier :**
  1. **searching the entire log is time-consuming**
  2. **we might unnecessarily redo transactions which have already output their updates to the database.**





# Checkpoints

- Streamline recovery procedure by periodically performing **check pointing**

-  1. Output all log records currently residing in main memory onto stable storage.
-  2. Output all modified buffer blocks to the disk.
-  3. Write a log record <checkpoint> onto stable storage.







# Checkpoints

- During recovery we need to consider only the most recent transaction  $T_i$  that started before the checkpoint, and transactions that started after  $T_i$ .
1. Scan backwards from end of log to find the most recent <checkpoint> record
  2. Continue scanning backwards till a record < $T_i$ start> is found.



# Checkpointing

assume transactions are run serially

3. Need only consider the part of log following above start record. Earlier part of log can be ignored during recovery, and can be erased whenever desired.
4. Scanning forward in the log, for all transactions starting from  $T_i$  or later with a  $\langle T_i, \text{commit} \rangle$ , execute  $\text{redo}(T_i)$ .
5. For all transactions (starting from  $T_i$  or later) with no  $\langle T_i, \text{commit} \rangle$ , execute  $\text{undo}(T_i)$ . (Done only in case of immediate modification.)



# Concurrent Executions<sup>Cont.</sup>

- **Concurrency control schemes** – mechanisms to achieve isolation
  - that is, to control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database






# Schedules

- **Schedule** – a sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed

- ⌘ a schedule for a set of transactions must consist of all instructions of those transactions
- ⌘ must preserve the order in which the instructions appear in each individual transaction.



# Schedule 1


- Let  $T_1$  transfer \$50 from  $A$  to  $B$ , and  $T_2$  transfer 10% of the balance from  $A$  to  $B$ .
- A **serial** schedule in which  $T_1$  is followed by  $T_2$ :

$T_1$	$T_2$
read( $A$ ) $A := A - 50$ write( $A$ ) read( $B$ ) $B := B + 50$ write( $B$ )	read( $A$ ) $temp := A * 0.1$ $A := A - temp$ write( $A$ ) read( $B$ ) $B := B + temp$ write( $B$ )

# Schedule 2

- A serial schedule where  $T_2$  is followed by  $T_1$

$T_1$	$T_2$
read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B)	read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B)






# Schedule 3

- Let  $T_1$  and  $T_2$  be the transactions defined previously. The following schedule is not a serial schedule, but it is *equivalent* to Schedule 1.

In Schedules 1, 2 and 3, the sum  $A + B$  is preserved.

$T_1$	$T_2$
read(A) $A := A - 50$ write(A)	read(A) $temp := A * 0.1$ $A := A - temp$ write(A)
read(B) $B := B + 50$ write(B)	read(B) $B := B + temp$ write(B)



# Schedule 4

- The following concurrent schedule does not preserve the value of  $(A + B)$ .

$T_1$	$T_2$
read( $A$ ) $A := A - 50$	read( $A$ ) $temp := A * 0.1$ $A := A - temp$ write( $A$ ) read( $B$ )
write( $A$ ) read( $B$ ) $B := B + 50$ write( $B$ )	$B := B + temp$ write( $B$ )

A: 150

temp: 20

B: 300



# Serializability

- **Basic Assumption** – Each transaction preserves database consistency.
- Thus serial execution of a set of transactions preserves database consistency.
- A (possibly concurrent) schedule is **serializable** if it is equivalent to a serial schedule.





# Serializability<sup>Cont.</sup>

- Different forms of schedule equivalence give rise to the notions:
  1. conflict serializability
  2. view serializability





# Serializability<sup>Cont.</sup>

- *Simplified view of transactions*
  - We ignore operations other than **read** and **write** instructions
  - We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes.
  - Our simplified schedules consist of only **read** and **write** instructions.



# Conflicting Instructions

- Instructions  $I_i$  and  $I_j$  of transactions  $T_i$  and  $T_j$  respectively, **conflict** if and only if there exists some item  $Q$  accessed by both  $I_i$  and  $I_j$ , and at least one of these instructions wrote  $Q$ .
  1.  $I_i = \text{read}(Q)$ ,  $I_j = \text{read}(Q)$ .  $I_i$  and  $I_j$  don't conflict.
  2.  $I_i = \text{read}(Q)$ ,  $I_j = \text{write}(Q)$ . They conflict.
  3.  $I_i = \text{write}(Q)$ ,  $I_j = \text{read}(Q)$ . They conflict
  4.  $I_i = \text{write}(Q)$ ,  $I_j = \text{write}(Q)$ . They conflict

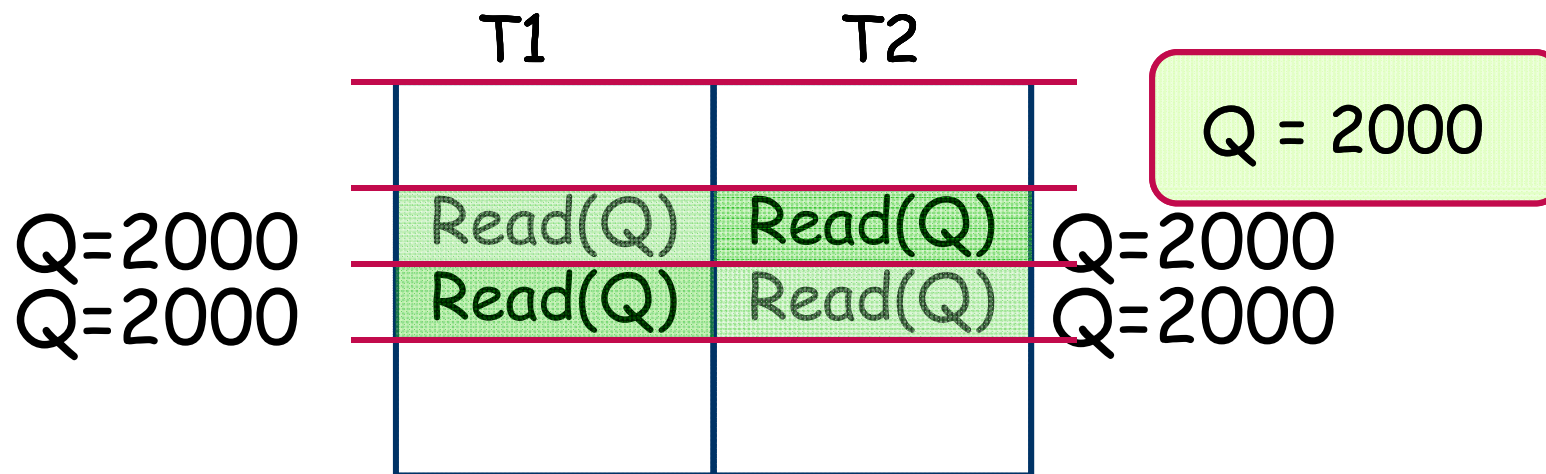




Cont.


# Conflicting Instructions

- If  $I_i$  and  $I_j$  are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.




# Conflict Serializability

- If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, we say that  $S$  and  $S'$  are **conflict equivalent**.
- We say that a schedule  $S$  is **conflict serializable** if it is conflict equivalent to a serial schedule



# Conflict Serializability<sup>Cont.</sup>

$T_1$	$T_2$
read(A) write(A)	read(A) write(A)
read(B) write(B)	
	read(B) write(B)



$T_1$	$T_2$
read(A) write(A) read(B) write(B)	read(A) write(A)
	read(B) write(B)

Schedule 3

Schedule 6

Conflict Serializable




# Conflict Serializability

- Example of a schedule that is not conflict serializable:

$T_3$	$T_4$
read( $Q$ )	write( $Q$ )
write( $Q$ )	

We are unable to swap instructions in the above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$ .



# View Serializability

- Let  $S$  and  $S'$  be two schedules with the same set of transactions.  $S$  and  $S'$  are **view equivalent** if the following three conditions are met, for each data item  $Q$ ,
  1. If in schedule  $S$ , transaction  $T_i$  reads the initial value of  $Q$ , then in schedule  $S'$  also transaction  $T_i$  must read the initial value of  $Q$ .



Cont.

# View Serializability

2. If in schedule  $S$  transaction  $T_i$  executes  $\text{read}(Q)$ , and that value was produced by transaction  $T_j$  (if any), then in schedule  $S'$  also transaction  $T_i$  must read the value of  $Q$  that was produced by the same  $\text{write}(Q)$  operation of transaction  $T_j$ .
3. The transaction (if any) that performs the final  $\text{write}(Q)$  operation in schedule  $S$  must also perform the final  $\text{write}(Q)$  operation in schedule  $S'$ .



# View Serializability

Cont.

- A schedule  $S$  is **view serializable** if it is view equivalent to a serial schedule.
- Every conflict serializable schedule is also view serializable.
- Below is a schedule which is view-serializable but *not* conflict serializable.

$T_3$	$T_4$	$T_6$
read( $Q$ )	write( $Q$ )	
write( $Q$ )		write( $Q$ )

$T_3$	$T_4$	$T_6$
read( $Q$ ) write( $Q$ )	write( $Q$ )	write( $Q$ )

Blind write




# Other Notions of Serializability

- The schedule below produces same outcome as the serial schedule  $\langle T_1, T_5 \rangle$ , yet is not conflict equivalent or view equivalent to it.

$T_1$	$T_5$
read(A) $A := A - 50$ write(A)	
	read(B) $B := B - 10$ write(B)
read(B) $B := B + 50$ write(B)	
	read(A) $A := A + 10$ write(A)

Determining such equivalence requires analysis of operations other than read and write.



# Recoverable Schedules

Need to address the effect of transaction failures on concurrently running transactions.

- **Recoverable schedule** — if a transaction  $T_j$  reads a data item previously written by a transaction  $T_i$ , then the commit operation of  $T_i$  appears before the commit operation of  $T_j$ .



# Recoverable Schedules

- The following schedule is not recoverable if  $T_9$  commits immediately after the read

$T_8$	$T_9$
read(A)	
write(A)	
fail	read(A)
read(B)	

commit

If  $T_8$  should abort,  $T_9$  would have read (and possibly shown to the user) an inconsistent database state. Hence, database must ensure that schedules are recoverable.

# Cascading Rollbacks

- **Cascading rollback** – a single transaction failure leads to a series of transaction rollbacks.
- Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable)

$T_{10}$	$T_{11}$	$T_{12}$
read(A) read(B) write(A)	read(A) write(A)	read(A)

If  $T_{10}$  fails,  $T_{11}$  and  $T_{12}$  must also be rolled back.

Can lead to the undoing of a significant amount of work



# Cascadeless Schedules

- **Cascadeless schedules** — cascading rollbacks cannot occur; for each pair of transactions  $T_i$  and  $T_j$  such that  $T_j$  reads a data item previously written by  $T_i$ , the commit operation of  $T_i$  appears before the read operation of  $T_j$ .
- Every cascadeless schedule is also recoverable
- It is desirable to restrict the schedules to those that are cascadeless




# Concurrency Control

- A database must provide a mechanism that will ensure that all possible schedules are
  - either conflict or view serializable,
  - and are recoverable
  - and preferably cascadeless



# Lock-Based Protocols

- A lock is a mechanism to control concurrent access to a data item
- Data items can be locked in two modes :
  - *exclusive (X) mode*. Data item can be both read as well as written
  - *shared (S) mode*. Data item can only be read.
- Lock requests are made to concurrency-control manager. Transaction can proceed only after lock is granted.





# Lock-Based Protocols

- Lock-compatibility matrix

	S	X
S	true	false
X	false	false


- A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions.



Cont.

# Lock-Based Protocols

- Any number of transactions can hold shared locks on an item,
  - but if any transaction holds an exclusive on the item no other transaction may hold any lock on the item
- If a lock cannot be granted, the requesting transaction is made to wait till all incompatible locks held by other transactions have been released. The lock is then granted.



# Pitfalls of Lock-Based Protocols

$T_3$	$T_4$
lock-X(B) read(B) $B := B - 50$ write(B)	
	lock-S(A) read(A) lock-S(B)
lock-X(A)	

Wait

T4 to release  
lock on A

Neither  $T_3$  nor  $T_4$   
can make progress  
Called **deadlock**

**Solution:**

Rollback one of  
the transactions

Wait

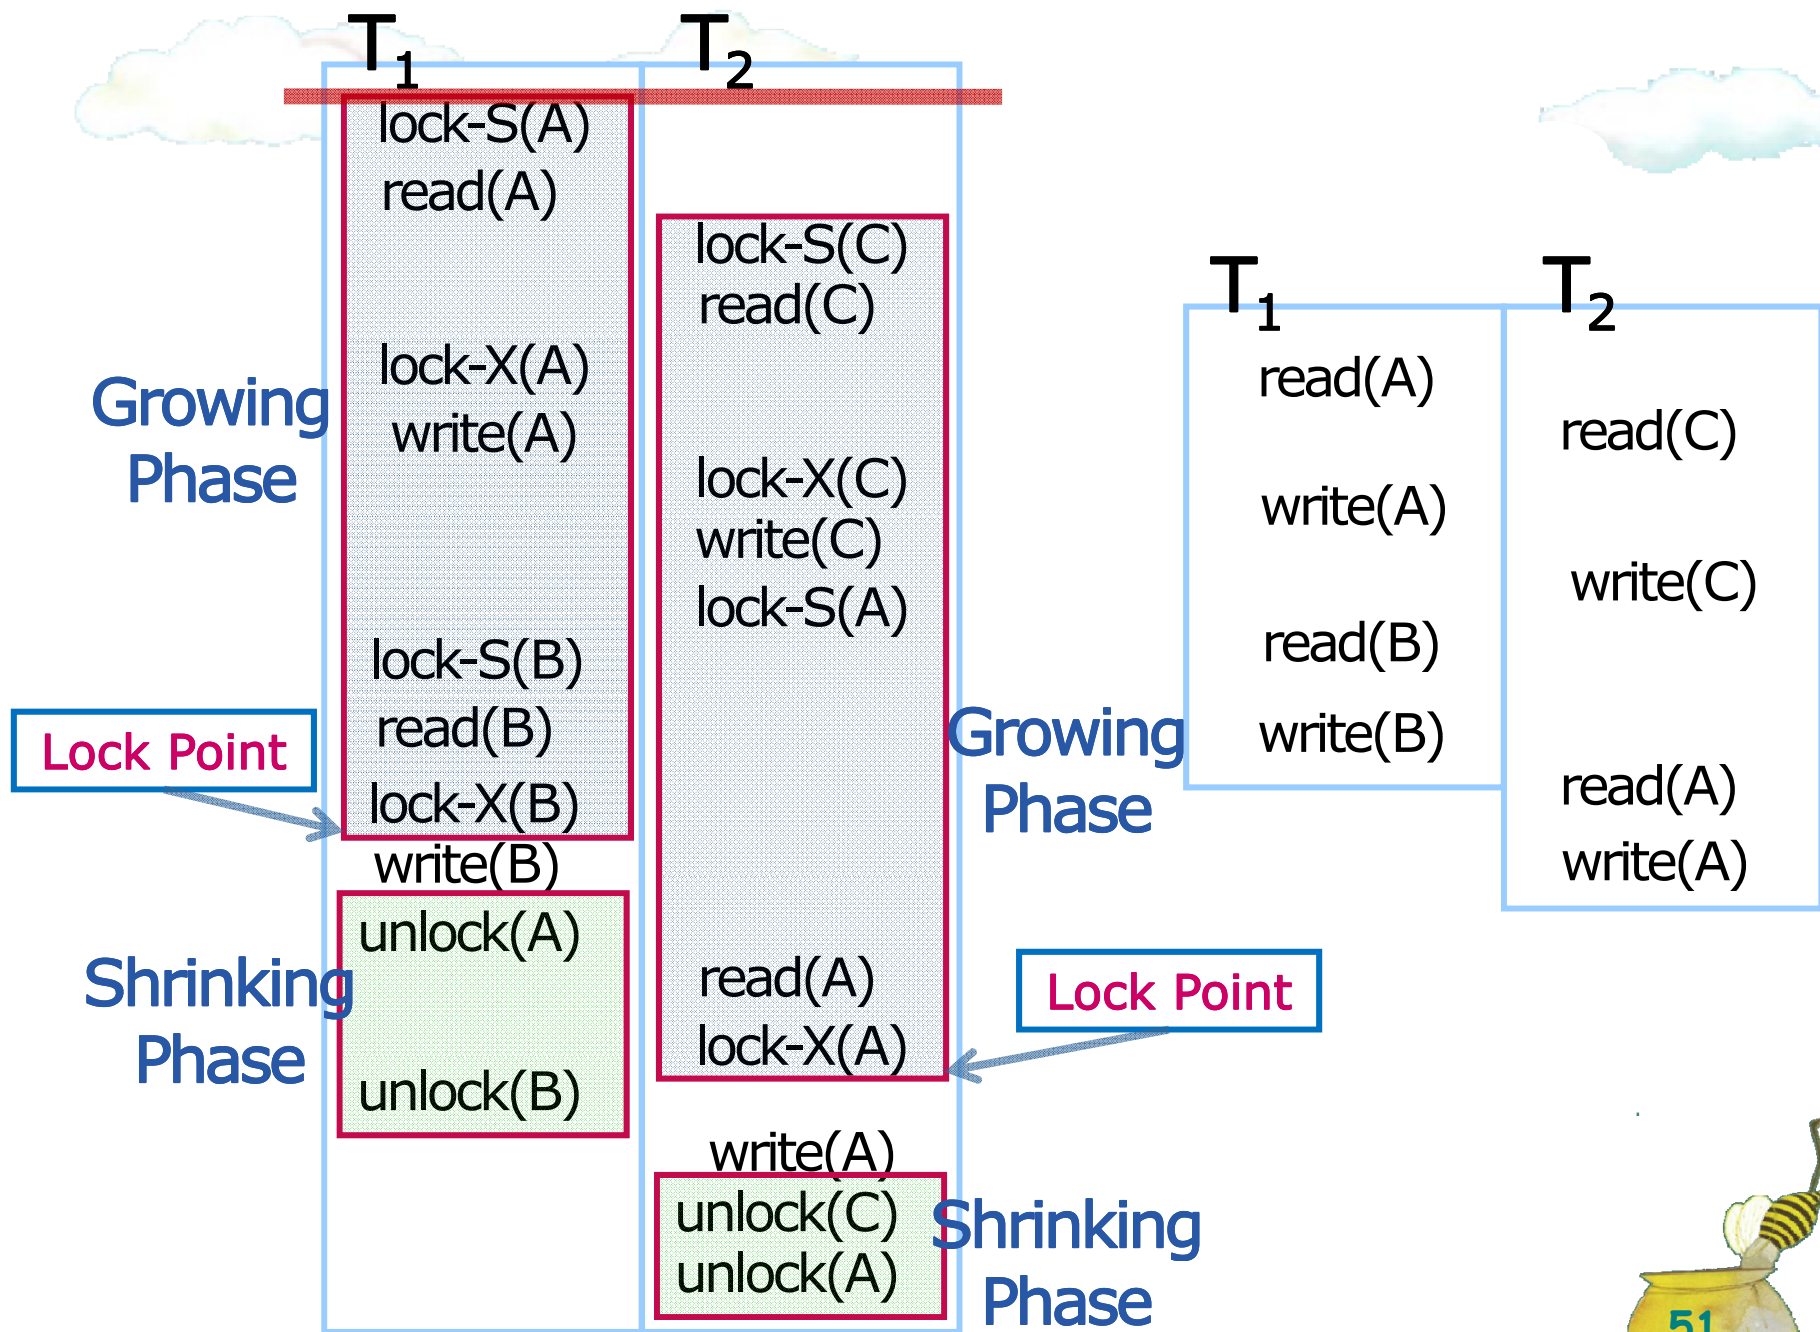
T3 to release  
lock on B



# Two-Phase Locking Protocol

- This is a protocol which ensures conflict-serializable schedules.
- **Phase 1: Growing Phase**
  - transaction may obtain locks
  - transaction may not release locks
- **Phase 2: Shrinking Phase**
  - transaction may release locks
  - transaction may not obtain locks





# Two-Phase Locking Protocol

- The protocol assures serializability. It can be proved that the transactions can be serialized in the order of their **lock points** (i.e. the point where a transaction acquired its final lock).
- Two-phase locking *does not* ensure freedom from deadlocks
- Cascading roll-back is possible





# Conclusions

